

# CA1

December 1, 2021

## 1 DAT300 - Compulsory assignment 1

### 2 Introduction

This notebook provides the solution for compulsory assignment 1 for Dat 300 (Deep learning) course. We were provided with a dataset from the titanic disaster. The dataset was based on 11 columns. Out of the 11 columns there is one target column as well that tells whether the person survived the accident or not. Our task was to train a model that predicted the target variable showing whether a person survived the event or not.

```
[1]: ## Import libraries
import pandas as pd
import numpy as np
import math

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.wrappers.scikit_learn import KerasClassifier

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
```

### 3 Data pre-processing and visualisation

```
[2]: # Load dataset from the train file
data = pd.read_csv("train.csv")
```

```
[3]: # Observe the dataframe
data.head()
```

```
[3]: Unnamed: 0    col_0    col_1    col_2    col_3    col_4 \
0          0  8.978818  7933.245770  18.006690 -0.001341  2.194868
1          1  3.994005   54.271828  36.536885 -0.000330  1.384795
2          2  0.997600   2.711767  32.193959  1.002930 -0.002403
3          3  8.985727  7988.253415  21.677670 -0.002819  2.195637
4          4  9.000558  8107.606049  65.403913  0.001092  2.197287

      col_5    col_6    col_7    col_8  target
0  9.130840  10.482944  0.074981  2.703268e+04    0.0
1  23.773308   6.993599 -0.010534  7.544705e+05    0.0
2  33.363202   7.006054  0.082351  1.153677e+06    0.0
3  10.488081  10.481805 -0.230482  5.169128e+04    0.0
4  12.416446  10.504148  0.087988  6.594806e+05    0.0
```

```
[4]: # Drop the unnamed column which is basically a serial number
data = data.drop(data.columns[0], axis=1)
```

```
[5]: # Check for null values
print(data.isna().sum())
```

```
col_0    0
col_1    0
col_2    0
col_3    0
col_4    0
col_5    0
col_6    0
col_7    0
col_8    0
target    0
dtype: int64
```

```
[6]: # Length of data
len(data)
```

```
[6]: 160000
```

## Visualization

```
[7]: data.describe()
```

```
[7]: count    160000.000000  160000.000000  160000.000000  160000.000000 \
mean         4.692211    2556.882075    38.842503    0.402870
std          3.160399    3731.293204    18.239105    0.814047
```

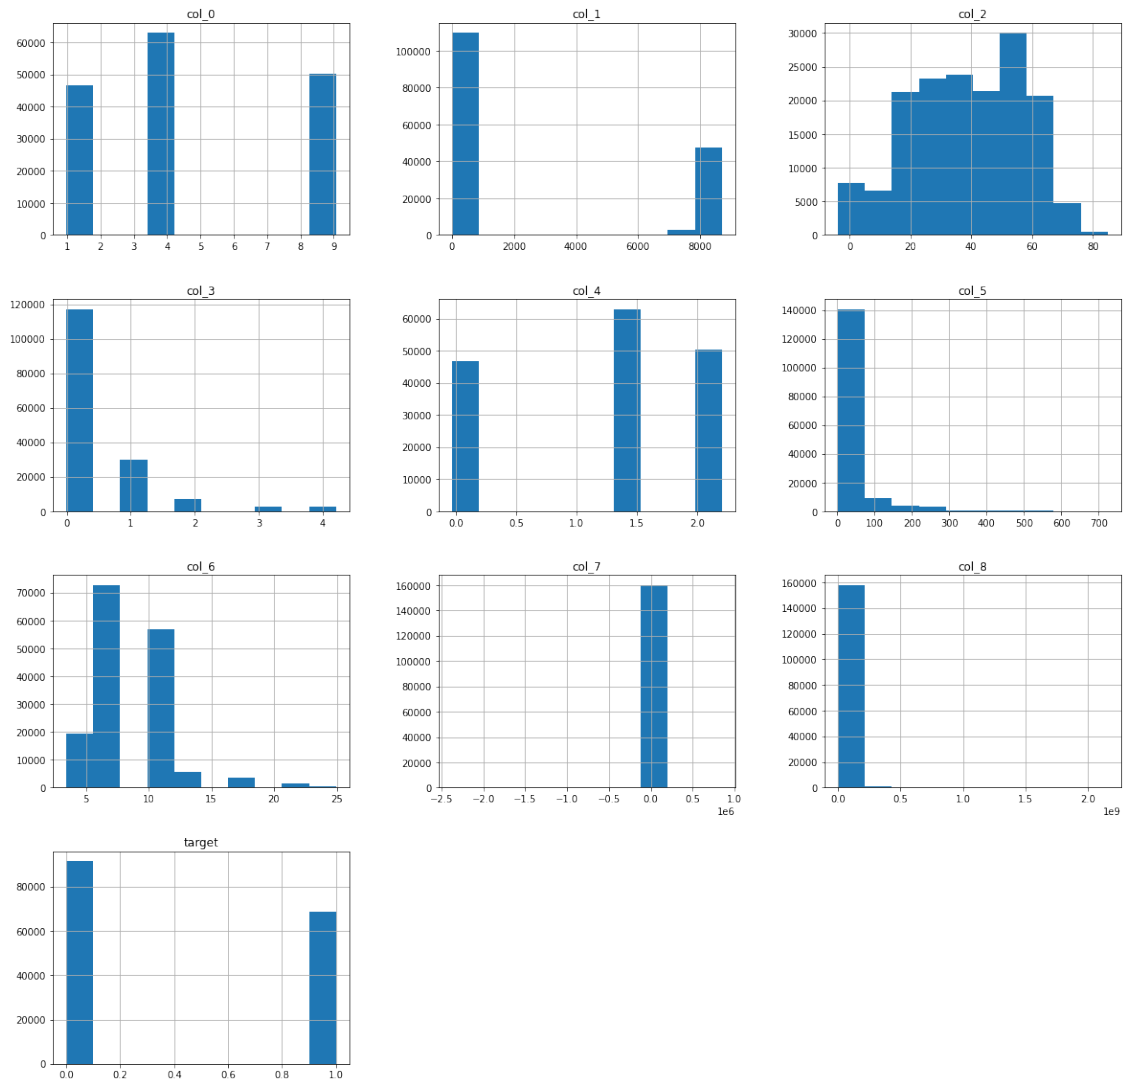
min	0.968253	2.633340	-3.826358	-0.009382
25%	1.009518	2.744277	24.172138	-0.001803
50%	3.998660	54.525056	39.534776	-0.000072
75%	8.980269	7944.772764	54.297277	0.994750
max	9.071419	8702.963926	84.941718	4.210325

	col_4	col_5	col_6	col_7 \
count	160000.000000	160000.000000	160000.000000	1.600000e+05
mean	1.235912	41.422999	8.485864	-1.722489e+01
std	0.862828	67.845301	3.249466	8.041504e+03
min	-0.032262	1.867981	3.432973	-2.373622e+06
25%	0.009473	11.472291	6.988763	-1.749591e-01
50%	1.385959	23.492842	7.011010	1.291813e-02
75%	2.195030	31.776878	10.493763	2.489355e-01
max	2.205129	724.802989	24.949926	8.504923e+05

	col_8	target
count	1.600000e+05	160000.000000
mean	1.271844e+07	0.428963
std	6.642855e+07	0.494929
min	6.502112e-03	0.000000
25%	1.010675e+05	0.000000
50%	4.643048e+05	0.000000
75%	2.488149e+06	1.000000
max	2.157407e+09	1.000000

If we observe this table we can see that columns (col\_5, col\_7 and col\_8) are showing significant differences. We might have to remove the outliers for these

```
[8]: # Histograms
data.hist(figsize=(20,20))
plt.show()
```

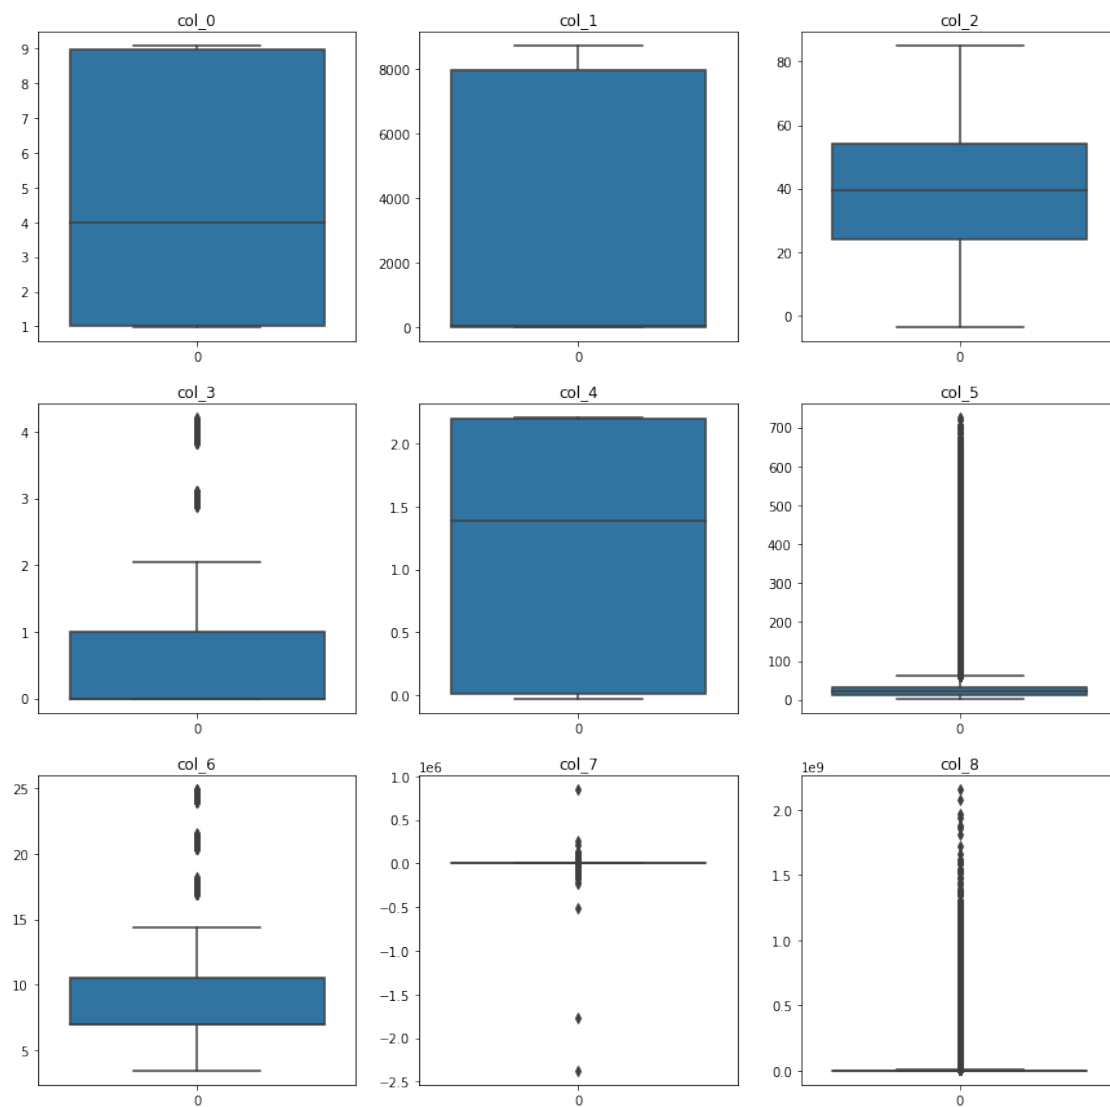


```
[9]: # Box Plots of the Columns
fig, axs = plt.subplots(3, 3, figsize=(15,15))
row = 0
col = 0
for x in axs:
    row_position = 0
    if row == 0:
        col = 0
    else:
        col += 1
    x[row_position].title.set_text(data.columns[col])
    sns.boxplot(data=data[data.columns[col]], ax=x[0])
    col += 1
    row_position += 1
```

```

x[row_position].title.set_text(data.columns[col])
sns.boxplot(data=data[data.columns[col]], ax=x[1])
col += 1
row_position += 1
x[row_position].title.set_text(data.columns[col])
sns.boxplot(data=data[data.columns[col]], ax=x[2])
row += 1

```



```


[10]: # Violen Plots of the Columns
fig, axs = plt.subplots(3, 3, figsize=(15,15))
row = 0
col = 0
for x in axs:

```

```

row_position = 0
if row == 0:
    col = 0
else:
    col += 1
x[row_position].title.set_text(data.columns[col])
sns.violinplot(data=data[data.columns[col]], ax=x[0])
col += 1
row_position += 1
x[row_position].title.set_text(data.columns[col])
sns.violinplot(data=data[data.columns[col]], ax=x[1])
col += 1
row_position += 1
x[row_position].title.set_text(data.columns[col])
sns.violinplot(data=data[data.columns[col]], ax=x[2])
row += 1


```



By observing the boxplots and violin plots, we confirm our theory of outliers in columns.

```
[11]: # Correlation matrix
plt.figure(figsize=(8,8))
sns.heatmap(data.corr(),annot=True)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f88b9118b10>
```



By observing the correlation heat map we see that col\_1 and col\_0 are highly correlated. Furthermore, col\_7 shows almost no correlation with the target variable. So we will remove these columns.

After removing the columns we will use the data to train machine learning models. We will not use the data with removed columns for training our ANN because the correlations could be non linear. We will let our ANN decide and train on its own.

```
[8]: data_ML = data.drop(columns=["col_0","col_7"])
```

```
[9]: data_ML.head()
```




```
[9]:
```

	col_1	col_2	col_3	col_4	col_5	col_6	\
0	7933.245770	18.006690	-0.001341	2.194868	9.130840	10.482944	
1	54.271828	36.536885	-0.000330	1.384795	23.773308	6.993599	
2	2.711767	32.193959	1.002930	-0.002403	33.363202	7.006054	
3	7988.253415	21.677670	-0.002819	2.195637	10.488081	10.481805	
4	8107.606049	65.403913	0.001092	2.197287	12.416446	10.504148	

	col_8	target
0	2.703268e+04	0.0
1	7.544705e+05	0.0
2	1.153677e+06	0.0
3	5.169128e+04	0.0
4	6.594806e+05	0.0

```
[10]: plt.figure(figsize=(8,8))
sns.heatmap(data_ML.corr(),annot=True)
```

```
[10]: <AxesSubplot:>
```



```
[11]: # Updating outliers and replacing them with median value
for x in data_ML.columns:
    if x == 'col_5' or x == 'col_7' or x == 'col_8':
        data_ML[x] = np.where(data_ML[x] > data_ML[x].quantile(0.90),
                               data_ML[x].quantile(0.50),
                               data_ML[x])
```

```
[16]: ## Box Plots of the Columns after outlier removal
# fig, axs = plt.subplots(3, 3, figsize=(15,15))
# row = 0
# col = 0
# for x in axs:
```

```

#     row_position = 0
#     if row == 0:
#         col = 0
#     else:
#         col += 1
#     x[row_position].title.set_text(data_ML.columns[col])
#     sns.boxplot(data=data_ML[data_ML.columns[col]], ax=x[0])
#     col += 1
#     row_position += 1
#     x[row_position].title.set_text(data_ML.columns[col])
#     sns.boxplot(data=data_ML[data_ML.columns[col]], ax=x[1])
#     col += 1
#     row_position += 1
#     x[row_position].title.set_text(data_ML.columns[col])
#     sns.boxplot(data=data_ML[data_ML.columns[col]], ax=x[2])
#     row += 1

```

```

[12]: # Dividing dataset into X(train) and y(test)
X_ML = data_ML.iloc[:, 0:7]
y_ML = data_ML.iloc[:,7]

```

```

[13]: # Test Train Split
X_train, X_test, y_train, y_test = train_test_split(X_ML, y_ML, test_size=0.2,
↳random_state=3, stratify=y_ML)

```

```

[14]: # Standarization
sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)

```

## 4 Methods

The Classification algorithm had a significant running time, especially the support vector machine. A grid search was tried in this case, but was not used because of the long running time. Random forest proved to be a good classifier for this, as well as the SVM. With more optimization, the latter could have had a better performance.

Setting outlier values to median value was done, to avoid exaggerated “Down-scaling” of data of lower value when scaled. This seemed to have a positive impact to the accuracy. Not all outlier values were replaced with median, as more than the upper 10% quantile did not contain all outliers for the columns. Therefore, the down-scaling could still be a factor. The downside to replacing outliers with median values is that important patterns can be removed, making it harder to perform a good prediction.

Column 7 were removed due to low correlation with the target. Column 0 were removed due to high correlation with column 1.

As for ANN training, different activation functions than the sigmoid activation function was tested in the last layer, but we experienced a drop in the accuracy. This may indicate that there are non-linear trends in the dataset.

Grid search was performed for parameter optimization, with good results.

For the ANN, lower amounts of epochs were tried, but with little success.

#### 4.0.1 Machine Learning models

```
[16]: # Support Vector Classifier
svc = SVC(gamma='auto')
svc.fit(X_train_std, y_train)
print('Support Vector training accuracy: ' , svc.score(X_train_std, y_train))
svc.predict(X_test_std)
print('Support Vector validation accuracy: ' , svc.score(X_test_std, y_test))
```

Support Vector training accuracy: 0.7955078125

Support Vector validation accuracy: 0.79525

```
[22]: # Grid search CV for Random Forrest
param_grid = [{'n_estimators': [100,200,300], 'criterion': ["gini",
↪ "entropy"]}]]
grid = GridSearchCV(estimator =
↪ RandomForestClassifier(random_state=1,n_jobs=-1), param_grid = param_grid)
grid.fit(X_train_std,y_train)
print('Random forrest with best params accuracy: ' , grid.best_score_)
print('Random forrest best params: ' ,grid.best_params_)
classifier_random = grid.best_estimator_
print('Random Forrest validation accuracy: ' , classifier_random.
↪ score(X_test_std, y_test))
```

Random forrest with best params accuracy: 0.7904609375

Random forrest best params: {'criterion': 'entropy', 'n\_estimators': 300}

Random Forrest validation accuracy: 0.79315625

#### 4.0.2 ANN Models

```
[7]: # Dividing dataset into X(train) and y(test)
X = data.iloc[:, 0:9]
y = data.iloc[:, 9]
# Test Train Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=3, stratify=y)
# Standarization
sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)
```

```
[18]: from tensorflow.keras import models
from tensorflow.keras import layers
from tensorflow.keras import optimizers
model_simple = models.Sequential([
    layers.Dense(128, activation='relu'),
    layers.Dense(128, activation='relu'),
    layers.Dense(4, activation='relu'),
    layers.Dense(1, activation='sigmoid')])
```

```
[19]: model_simple.compile(optimizer='rmsprop',
    loss='binary_crossentropy',
    metrics=['accuracy'])
```

```
[21]: model_simple.fit(X_train_std,
    y_train,
    epochs=30,
    batch_size=512,
    validation_data=(X_test_std, y_test))
```

```
Epoch 1/30
250/250 [=====] - 1s 4ms/step - loss: 0.4348 -
accuracy: 0.8007 - val_loss: 0.4340 - val_accuracy: 0.8001
Epoch 2/30
250/250 [=====] - 1s 4ms/step - loss: 0.4345 -
accuracy: 0.8005 - val_loss: 0.4360 - val_accuracy: 0.7999
Epoch 3/30
250/250 [=====] - 1s 4ms/step - loss: 0.4349 -
accuracy: 0.8005 - val_loss: 0.4340 - val_accuracy: 0.7997
Epoch 4/30
250/250 [=====] - 1s 4ms/step - loss: 0.4342 -
accuracy: 0.8013 - val_loss: 0.4330 - val_accuracy: 0.8010
Epoch 5/30
250/250 [=====] - 1s 4ms/step - loss: 0.4338 -
accuracy: 0.8010 - val_loss: 0.4323 - val_accuracy: 0.7996
Epoch 6/30
```

250/250 [=====] - 1s 4ms/step - loss: 0.4337 -  
accuracy: 0.8016 - val\_loss: 0.4308 - val\_accuracy: 0.8009  
Epoch 7/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4331 -  
accuracy: 0.8012 - val\_loss: 0.4327 - val\_accuracy: 0.8016  
Epoch 8/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4330 -  
accuracy: 0.8016 - val\_loss: 0.4305 - val\_accuracy: 0.8012  
Epoch 9/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4326 -  
accuracy: 0.8016 - val\_loss: 0.4313 - val\_accuracy: 0.8008  
Epoch 10/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4321 -  
accuracy: 0.8015 - val\_loss: 0.4313 - val\_accuracy: 0.8012  
Epoch 11/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4314 -  
accuracy: 0.8022 - val\_loss: 0.4312 - val\_accuracy: 0.7991  
Epoch 12/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4309 -  
accuracy: 0.8023 - val\_loss: 0.4315 - val\_accuracy: 0.8011  
Epoch 13/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4311 -  
accuracy: 0.8033 - val\_loss: 0.4296 - val\_accuracy: 0.8019  
Epoch 14/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4301 -  
accuracy: 0.8039 - val\_loss: 0.4278 - val\_accuracy: 0.8027  
Epoch 15/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4293 -  
accuracy: 0.8033 - val\_loss: 0.4282 - val\_accuracy: 0.8009  
Epoch 16/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4291 -  
accuracy: 0.8040 - val\_loss: 0.4260 - val\_accuracy: 0.8020  
Epoch 17/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4286 -  
accuracy: 0.8045 - val\_loss: 0.4299 - val\_accuracy: 0.8023  
Epoch 18/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4277 -  
accuracy: 0.8043 - val\_loss: 0.4253 - val\_accuracy: 0.8038  
Epoch 19/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4275 -  
accuracy: 0.8050 - val\_loss: 0.4260 - val\_accuracy: 0.8040  
Epoch 20/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4272 -  
accuracy: 0.8055 - val\_loss: 0.4273 - val\_accuracy: 0.8042  
Epoch 21/30  
250/250 [=====] - 1s 4ms/step - loss: 0.4264 -  
accuracy: 0.8056 - val\_loss: 0.4257 - val\_accuracy: 0.8041  
Epoch 22/30

```

250/250 [=====] - 1s 4ms/step - loss: 0.4262 -
accuracy: 0.8063 - val_loss: 0.4299 - val_accuracy: 0.8044
Epoch 23/30
250/250 [=====] - 1s 4ms/step - loss: 0.4257 -
accuracy: 0.8058 - val_loss: 0.4283 - val_accuracy: 0.8048
Epoch 24/30
250/250 [=====] - 1s 4ms/step - loss: 0.4252 -
accuracy: 0.8067 - val_loss: 0.4261 - val_accuracy: 0.8045
Epoch 25/30
250/250 [=====] - 1s 4ms/step - loss: 0.4251 -
accuracy: 0.8063 - val_loss: 0.4298 - val_accuracy: 0.8047
Epoch 26/30
250/250 [=====] - 1s 4ms/step - loss: 0.4244 -
accuracy: 0.8072 - val_loss: 0.4269 - val_accuracy: 0.8045
Epoch 27/30
250/250 [=====] - 1s 4ms/step - loss: 0.4244 -
accuracy: 0.8073 - val_loss: 0.4251 - val_accuracy: 0.8057
Epoch 28/30
250/250 [=====] - 1s 4ms/step - loss: 0.4244 -
accuracy: 0.8076 - val_loss: 0.4236 - val_accuracy: 0.8052
Epoch 29/30
250/250 [=====] - 1s 4ms/step - loss: 0.4239 -
accuracy: 0.8077 - val_loss: 0.4317 - val_accuracy: 0.8033
Epoch 30/30
250/250 [=====] - 1s 4ms/step - loss: 0.4245 -
accuracy: 0.8080 - val_loss: 0.4251 - val_accuracy: 0.8062

```

[21]: <keras.callbacks.History at 0x7f5830ccba90>

Since the accuracy is pretty even from the lowest epoch values, and the loss is low, we had to perform a grid search to see which was the optimal amount of epochs.

[27]: `model_simple`

[27]: <keras.engine.sequential.Sequential at 0x7f88b727efd0>

### Grid Search for keras ANN

```

[8]: # Calculate number of nodes for each layer
def FindLayerNodesLinear(n_layers, first_layer_nodes, last_layer_nodes):
    layers = []

    nodes_increment = (last_layer_nodes - first_layer_nodes) / (n_layers-1)
    nodes = first_layer_nodes
    for i in range(1, n_layers+1):
        layers.append(math.ceil(nodes))
        nodes = nodes + nodes_increment

```

```
return layers
```

```
[9]: # Create model
def createmodel(n_layers, first_layer_nodes, last_layer_nodes, activation_func,
    ↪loss_func):
    model = Sequential()
    n_nodes = FindLayerNodesLinear(n_layers, first_layer_nodes,
    ↪last_layer_nodes)
    for i in range(1, n_layers):
        if i==1:
            model.add(Dense(first_layer_nodes, input_dim=X_train.shape[1],
    ↪activation=activation_func))
        else:
            model.add(Dense(n_nodes[i-1], activation=activation_func))

    #Finally, the output layer should have a single node in binary
    ↪classification
    model.add(Dense(1, activation=activation_func))
    model.compile(optimizer='rmsprop', loss=loss_func, metrics = ["accuracy"])
    ↪#note: metrics could also be 'mse'

    return model

##Wrap model into scikit-learn
model = KerasClassifier(build_fn=createmodel, verbose = False)
```

```
[10]: ## Define params for grid search
activation_funcs = ['sigmoid', 'relu']
loss_funcs = ['binary_crossentropy']
param_grid = dict(n_layers=[2,3,4], first_layer_nodes = [128,64,32,16],
    ↪last_layer_nodes = [4], activation_func = activation_funcs, loss_func =
    ↪loss_funcs, batch_size = [128,256,512,1024], epochs = [10,20,30])
grid = GridSearchCV(estimator = model, param_grid = param_grid)
```

```
[11]: # Train models
grid.fit(X_train_std,y_train)
```

```
[11]: GridSearchCV(cv=None, error_score=nan,
    estimator=<keras.wrappers.scikit_learn.KerasClassifier object at
    0x7f583d355e10>,
    iid='deprecated', n_jobs=None,
    param_grid={'activation_func': ['sigmoid', 'relu'],
    'batch_size': [128, 256, 512, 1024],
    'epochs': [10, 20, 30],
    'first_layer_nodes': [128, 64, 32, 16],
    'last_layer_nodes': [4],
```



```

        'loss_func': ['binary_crossentropy'],
        'n_layers': [2, 3, 4]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)

```

```

[12]: print('Grid Search on ANN (Best Score): ', grid.best_score_)
      print('Grid Search on ANN (Best Params): ', grid.best_params_)

```

```

Grid Search on ANN (Best Score): 0.7983437657356263
Grid Search on ANN (Best Params): {'activation_func': 'relu', 'batch_size':
256, 'epochs': 30, 'first_layer_nodes': 128, 'last_layer_nodes': 4, 'loss_func':
'binary_crossentropy', 'n_layers': 4}

```

```

[13]: best_ANN = grid.best_estimator_

```

```

[14]: def make_submission(model, name = 'submission.csv'):
      data_test_final = pd.read_csv("test.csv")
      data_test_final = data_test_final.iloc[:, 1:10]
      print (data_test_final.head())
      data_test_std = sc.transform(data_test_final)
      test_pred_final = model.predict(data_test_std)
      print (test_pred_final)
      # Convert probabilities to classes
      targets = (test_pred_final > 0.5).astype("int64").astype("float64")
      csv_data = {'Survived': targets[:, 0]}
      df = pd.DataFrame(csv_data)
      df.index.name = 'Id'
      print(df.head())
      df.to_csv(name)

```

```

[22]: make_submission(model_simple, 'submission_simple.csv')

```

	col_0	col_1	col_2	...	col_6	col_7	col_8
0	4.022784	55.856406	29.922403	...	7.014724	-0.083142	601273.608998
1	3.992970	54.215653	26.485782	...	6.985344	0.288968	206517.816317
2	8.992121	8039.491762	59.917066	...	10.496820	-0.039709	222066.359095
3	0.984910	2.677570	4.797783	...	7.005128	470.613913	22911.001098
4	1.000011	2.718311	32.485426	...	3.494827	0.859416	85149.762596

```

[5 rows x 9 columns]
[[0.7827109 ]
 [0.08352521]
 [0.09349078]
 ...
 [0.27765667]
 [0.43107867]
 [0.01479951]]

```

	Survived
Id	
0	1.0
1	0.0
2	0.0
3	1.0
4	0.0

```
[17]: make_submission(best_ANN, 'submission_grid.csv')
```

	col_0	col_1	col_2	...	col_6	col_7	col_8
0	4.022784	55.856406	29.922403	...	7.014724	-0.083142	601273.608998
1	3.992970	54.215653	26.485782	...	6.985344	0.288968	206517.816317
2	8.992121	8039.491762	59.917066	...	10.496820	-0.039709	222066.359095
3	0.984910	2.677570	4.797783	...	7.005128	470.613913	22911.001098
4	1.000011	2.718311	32.485426	...	3.494827	0.859416	85149.762596

[5 rows x 9 columns]

```
[[1.]
 [0.]
 [0.]
 ...
 [0.]
 [1.]
 [0.]]
```

	Survived
Id	
0	1.0
1	0.0
2	0.0
3	1.0
4	0.0

## 5 Results

## 6 Discussion / conclusion

Provide a summary of the assignment: (you are required to address **the first three** points of the list below) - obstacles / problems you have met regarding the modelling proces - degree of success - given more time, what would be done differently - further comments (if any)

- Challenges in the modelling was the large variance in the data; large amounts of data were grouped in “high value” outlier groups. These were too many to be removed, so the effect of this could be that the scaled data would leave the majority of values very small, making it harder for the algorithm to detect trends in the lower valued data, and apply the right weights here. A challenge with ANN is the vast amount of options to go with - leaving the likelihood to achieve the optimal solution smaller.

- The model achieved an accuracy of around 80%, which is quite good with a difficult data set. The ANN model did however not perform very much better than the simpler models, which is lower than one would expect with a more advanced model.
- A variety of options could be tested; amount layers, different layer sizes and various preprocessing options like feature engineering etc. We encountered large differences by making small alterations to the i.e. batch size and epochs of the model,